

## CS7070 Big Data Analytics

### Programming Project Assignment

#### Guidelines followed while handling data:

- All the words in the data file are converted into lower case and punctuations present at the leading and trailing of the words are stripped.  
For example: "We won't let you go!" => <we,1> <won't,1><let,1><you,1><go,1>
- The punctuation within the word is left so that the two actual meanings of the words "its" and "it's" is preserved.

**Note:** However, the results of the raw data are also presented in the **appendix** of this document.

**Note:** The code presented in this document is the modified code from the source:

<http://www.michael-noll.com/tutorials/writing-an-hadoop-mapreduce-program-in-python/>

1. Design and execute a MapReduce program to produce the frequencies of all the words in the book collection, retaining only those words whose frequencies are greater than 5. Submit the following:

- a. Your commented code for the Mapper and Reducer

**Mapper Code:** wc\_1\_1\_map.py

```
#!/usr/bin/env python
```

```
import sys
```

```
import string
```

```
def read_input(file):
```

```
    for line in file:
```

```
        # strip the leading and trailing spaces
```

```
        # And, split the line into words
```

```
        yield line.strip().split()
```

```
def main(separator='\t'):
```

```
    # input comes from STDIN (standard input)
```

```
    data = read_input(sys.stdin)
```

```
    for words in data:
```

```
        # Each word in words is converted into lower case
```

```
        # and leading an trailing punctuations are removed;
```

```
        # write each word tab-delimited with the trivial count 1 to STDOUT (standard output)
```

```
        # what we output here will be the input for the
```

```
        # Reduce step, i.e. the input for reducer.py
```

```
    for word in words:
```

```
        word = word.lower().strip(string.punctuation)
```

```
        if word:
```

```
            print '%s%s%d' % (word, separator, 1)
```

```
if __name__ == "__main__":
```

```
    main()
```

**Reducer Code:** wc\_1\_1\_red.py

```
#!/usr/bin/env python
"""A more advanced Reducer, using Python iterators and generators."""

from itertools import groupby
from operator import itemgetter
import sys

def read_mapper_output(file, separator='\t'):
    for line in file:
        yield line.rstrip().split(separator, 1)

def main(separator='\t'):
    # input comes from STDIN (standard input)
    data = read_mapper_output(sys.stdin, separator=separator)
    # groupby groups multiple word-count pairs by word,
    # and creates an iterator that returns consecutive keys and their group:
    # current_word - string containing a word (the key)
    # group - iterator yielding all ["<current_word>", "<count>"] items
    for current_word, group in groupby(data, itemgetter(0)):
        try:
            total_count = sum(int(count) for current_word, count in group)
            # write only words whose total count > 5
            if total_count > 5:
                print "%s%s%d" % (current_word, separator, total_count)
        except ValueError:
            # count was not a number, so silently discard this item
            pass

if __name__ == "__main__":
    main()
```

**Command** used to execute the above code in Hadoop:

```
$ hadoop jar /usr/hdp/3.0.0.0-1634/hadoop-mapreduce/hadoop-streaming.jar -file
wc_1_1_map.py -mapper wc_1_1_map.py -file wc_1_1_red.py -reducer wc_1_1_red.py -
input /user/kumarlt/books/input/* -output BDA/PPA_1/results/wc_1_1_cres
```

b. First 50 words and their frequencies from your program's output file.  
The first 50 words from the program output file are numericals as show in the screen-shot on the next page.

**Each word is tab delimited with its frequency.**

1	255
1.a	21
1.b	21
1.c	42
1.d	21
1.e	42
1.e.1	105
1.e.2	21
1.e.3	21
1.e.4	21
1.e.5	21
1.e.6	21
1.e.7	63
1.e.8	84
1.e.9	63
1.f	21
1.f.1	21
1.f.2	21
1.f.3	86
1.f.4	21
1.f.5	21
1.f.6	21
10	84
100	13
1000	7
103	6
104	7
105	8
106	6
107	6
11	44
110	6
12	44
120	8
13	36
14	34
15	72
1500	23
1502	7
1512	7
1513	8
15th	7
16	29
17	33
17-	7
18	42
18th	6
19	33
2	167
20	96

c. Last 50 words and their frequencies from your program's output file.

```
"ten      6
"thank    12
"that     166
"that's   41
"that's   10
"the      253
"then     40
"there    90
"there's
"there's
"these    15
"they     53
"this     69
"thou     19
"though   11
"to       54
"tom      38
>true     6
"two      8
"upon     12
"very     29
"was      15
"we       74
"well     307
"well,"   24
"were     7
"what     250
"what's   37
"what's   12
"when     48
"where    23
"which    14
"while    6
"who      46
"who's    8
"why      201
"why?"    7
"will     27
"with     22
"without
"would    9
"yes      183
"yes,"    28
"yes."    27
"yet      8
"you      352
"you'll   6
"you're   7
"your     52
```

2. Design and execute a MapReduce program to produce frequencies of all 2-gram word-pairs in the book collection, retaining only those 2-grams whose frequencies are greater than 5. Submit the following:

a. Your commented MapReduce program to produce the frequencies of all the 2-grams in all the books.

**Mapper Code:** wc\_2\_1\_map.py

```
#!/usr/bin/env python
```

```
import sys
```

```
import string
```

```
def read_input(file):
```

```
    for line in file:
```

```
        # split the line into words
```

```
        yield line.strip().split()
```

```
def main(separator='\t'):
```

```
    # input comes from STDIN (standard input)
```

```
    data = read_input(sys.stdin)
```

```
    # initializing lastword to empty string
```

```
    lastword = ""
```

```
    for words in data:
```

```
        # Each word in words is converted into lower case
```

```
        # and leading an trailing punctuations are removed;
```

```
        # write each word tab-delimited with the next word and again
```

```
        # tab-delimited with trivial count 1 to STDOUT (standard output)
```

```
        # what we output here will be the input for the
```

```
        # Reduce step, i.e. the input for reducer.py
```

```
    for i in range(len(words)-2):
```

```
        word_1 = words[i].lower().strip(string.punctuation)
```

```
        word_2 = words[i+1].lower().strip(string.punctuation)
```

```
        if word_1 and word_2:
```

```
            print '%s%s%s%s%d' % (word_1, separator, word_2,separator,1)
```

```
    # This is code below deals with producing the bigrams with the word
```

```
    # from current line last word and next line first word
```

```
    if lastword:
```

```
        print '%s%s%s%s%d' % (lastword, separator, words[0].lower()
```

```
        \.strip(string.punctuation),separator,1)
```

```
        lastword=words[-1].lower().strip(string.punctuation)
```

```
if __name__ == "__main__":
```

```
    main()
```

### Reducer Code: wc\_2\_1\_red.py

```
#!/usr/bin/env python
from itertools import groupby
from operator import itemgetter
import sys

def read_mapper_output(file, separator='\t'):
    for line in file:
        yield line.rstrip().split(separator, 2)

def main(separator='\t'):
    # input comes from STDIN (standard input)
    data = read_mapper_output(sys.stdin, separator=separator)
    # groupby groups multiple word-word-count pairs by firstword,
    # and creates an iterator that returns consecutive keys and their group:
    # current_firstword - string containing a word (the key)
    # group - iterator yielding all ["<current_firstword>",
    # "<current_secondword>", "<count>"] items
    for current_firstword, group_1 in groupby(data, itemgetter(0)):
        try:
            # groupby groups multiple word-word-count pairs by secondword,
            # and creates an iterator that returns consecutive keys and their group:
            # current_secondword - string containing a word (the key)
            # group - iterator yielding all ["<current_firstword>",
            # "<current_secondword>", "<count>"] items
            for current_secondword, group_2 in groupby(group_1, itemgetter(1)):
                try:
                    total_count = sum(int(count) for firstword, secondword, count in group_2)
                    if total_count > 5:
                        print "%s%s%s%s%d" % (current_firstword, separator, \
                            current_secondword, separator, total_count)
                except ValueError:
                    pass
        except ValueError:
            # count was not a number, so silently discard this item
            pass

if __name__ == "__main__":
    main()
```

**Command** used to execute the above code in Hadoop:

```
$ hadoop jar /usr/hdp/3.0.0.0-1634/hadoop-mapreduce/hadoop-streaming.jar -file
wc_2_1_map.py -mapper wc_2_1_map.py -file wc_2_1_red.py -reducer wc_2_1_red.py -
input /user/kumarl/books/input/* -output BDA/PPA_1/results/wc_2_1_cres
```

b. First 50 2-grams and their frequencies from your program's output file.

Each line contains bigram( two words delimited by tab )and its frequency.

```

1.a      by      20
1.d      the     20
1.e      unless  20
1.e.2    if      20
1.e.3    if      20
1.e.4    do      20
1.e.5    do      20
1.e.6    you     20
1.e.9    if      20
1.f.1    project 20
1.f.2    limited 20
1.f.4    except  20
1.f.5    some    20
1.f.6    indemnity 20
1500     west    17
2        information 6
20       of      7
2001     the     21
2008     ebook   9
4557     melan   21
5        general  7
5,000    are     20
50       states  6
50       states  9
60       days    12
64-6221541 its     13
64-6221541 its     6
801      596-1887 19
809      north   19
84116    801     21
90       days    11
90       days    7
99712    but     21
a-going  to      7
abide    by      14
able     to      15
able     to      8
able     to      9
able     to      11
able     to      10
able     to      53
able     to      62
able     to      36
able     to      7
able     to      11
able     to      13
abstain  from    6
abundance of     6
abuse    of      6

```

c. Last 50 2-grams and their frequencies from your program's output file.

```

white    rabbit    7
widest   variety   13
willing  to          8
window   and         6
wish     to         9
wish     to         8
wished   to         6
wished   i          6
wishing  to        19
wisht    i          6
wrap     in         6
wreath   of         6
  by     and         8
  by     and         6
  do     you         7
  it     was         6
  it     was         6
  pretty soon      12
'ah!'    said       10
'behave  well       6
'it      is         6
'my      mother    7
'of      course     6
'off     with       8
'we      will       6
'well,'  said        9
'whither away      7
'yes,'   said        6
"all     right      8
"are     you         7
"are     you         8
"blame   it          8
"defects," such     10
"do      you         6
"do      you         6
"have    you         7
"he      is          6
"information about   9
"it      is          6
"it      is          6
"miss    manette    8
"my      dear        7
"plain   vanilla    17
"project gutenber"   31
"queequeg," said     7
"she     is          6
"this    is          6
"upon    my          8
"well,"  said        6

```



3. Design and execute a MapReduce program to produce the top 100 most frequent words in the book collection. You may need two rounds of Map and Reduce processors. Submit the following:

a. Your commented code for the Mappers and Reducers.

I have used output file of the first question as input to the below mapper.

**Mapper Code:** wc\_3\_1\_map.py

```
#!/usr/bin/env python
```

```
import sys
```

```
def read_input(file,separator = '\t'):
```

```
    for line in file:
```

```
        # split the line into word and value
```

```
        yield line.strip().split(separator,1)
```

```
def main(separator='\t'):
```

```
    # input comes from STDIN (standard input)
```

```
    data = read_input(sys.stdin)
```

```
    for line in data:
```

```
        # collect word and value and reverse its role
```

```
        word = line[0]
```

```
        value= int(line[1])
```

```
        print '%d%s%s' % (value, separator, word)
```

```
if __name__ == "__main__":
```

```
    main()
```

**Reducer Code:** wc\_3\_1\_red.py

```
#!/usr/bin/env python
```

```
"""A more advanced Reducer, using Python iterators and generators."""
```

```
from itertools import groupby
```

```
from operator import itemgetter
```

```
import sys
```

```
import collections
```

```
def read_mapper_output(file, separator='\t'):
```

```
    for line in file:
```

```
        yield line.rstrip().split(separator, 1)
```

```
def main(separator='\t'):
```

```
    # input comes from STDIN (standard input)
```

```
    data = read_mapper_output(sys.stdin, separator=separator)
```

```
    # groupby groups multiple value-word pairs by value,
```

```
    # and creates an iterator that returns consecutive words and their group:
```

```

# current_word - string containing a word (the key)
# group - iterator yielding all ["<current_word>", "<count>"] items
# combine all the words in this group by separator tab
# and output the frequency-# of words - words group separated by tab
for current_value, group in groupby(data, itemgetter(0)):
    try:
        total_word_group = "\t".join(word for current_value, word in group)
        print "%d%s%d%s" % (int(current_value), separator, \
            int(len(total_word_group.split(separator))), separator, total_word_group)
    except ValueError:
        # count was not a number, so silently discard this item
        pass

if __name__ == "__main__":
    main()

```

**NOTE:** I have used the output of the first question solution as input to this above mapper and reducer. And also, Hadoop inbuilt key comparator options are set to sort decreasing order of the frequency.

**Command** used to execute the above code in Hadoop:

```

$ hadoop jar /usr/hdp/3.0.0.0-1634/hadoop-mapreduce/hadoop-streaming.jar -D
mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFieldBasedCompar
ator -D mapred.text.key.comparator.options=-k1,1nr -file wc_3_1_map.py -mapper
wc_3_1_map.py -file wc_3_1_red.py -reducer wc_3_1_red.py -input
BDA/PPA_1/results/wc_1_1_cres/part-00000 -output BDA/PPA_1/results/wc_3_1_cres

```

b. The list of 100 most frequent words and their frequencies in the book collection.

108508	1	the
63903	1	and
53703	1	of
49034	1	to
34427	1	a
29901	1	in
24368	1	i
23421	1	that
20848	1	he
20356	1	it
19261	1	his
18522	1	was
15010	1	is
14921	1	with
14009	1	for
13940	1	as
12736	1	you
12432	1	but
11851	1	not
10809	1	be
10226	1	had
9887	1	by
9793	1	on
9524	1	at
9378	1	all
8950	1	him
8615	1	my
8595	1	her
8435	1	this
8289	1	they
8068	1	from
8049	1	have
7890	1	or
7858	1	so
7274	1	which
6972	1	she
6770	1	me
5997	1	there
5870	1	when
5800	1	said
5789	1	no
5722	1	are
5713	1	their
5711	1	one
5397	1	we
5373	1	were
5167	1	if
5160	1	them
5047	1	then

4506	1	what	
4361	1	out	
4212	1	an	
4128	1	would	
3969	1	up	
3853	1	will	
3773	1	been	
3717	1	any	
3552	1	now	
3533	1	some	
3532	1	more	
3522	1	who	
3511	1	man	
3474	1	do	
3465	1	could	
3346	1	your	
3324	1	into	
3123	1	other	
2947	1	time	
2923	1	such	
2857	1	very	
2717	1	upon	
2711	1	may	
2677	1	see	
2594	2	like	can
2574	1	down	
2561	1	before	
2549	1	our	
2545	1	shall	
2514	1	than	
2429	1	about	
2427	1	little	
2382	1	must	
2332	1	has	
2300	1	did	
2229	2	over	only
2195	1	know	
2182	1	mr	
2133	1	these	
2125	1	should	
2123	2	where	men
2107	1	great	
2046	1	again	
2020	1	come	
1989	1	good	

The above output represents the frequency-tab-#of words of that frequency-tab-words list.  
These are the top 100 words.

By:  
Varun Raj Rayabarapu

## Appendix:

This appendix consists solutions to the above questions when the raw data is handled.

### 1. a. Code:

**Mapper Code:** wc\_1\_1\_map.py

```
#!/usr/bin/env python
import sys
def read_input(file):
    for line in file:
        # split the line into words
        yield line.strip().split()
def main(separator='\t'):
    # input comes from STDIN (standard input)
    data = read_input(sys.stdin)
    for words in data:
        # write the results to STDOUT (standard output);
        # what we output here will be the input for the
        # Reduce step, i.e. the input for reducer.py
        # tab-delimited; the trivial word count is 1
        for word in words:
            print '%s%s%d' % (word, separator, 1)
if __name__ == "__main__":
    main()
```

**Reducer Code:** wc\_1\_1\_red.py

```
#!/usr/bin/env python
from itertools import groupby
from operator import itemgetter
import sys
def read_mapper_output(file, separator='\t'):
    for line in file:
        yield line.rstrip().split(separator, 1)
def main(separator='\t'):
    # input comes from STDIN (standard input)
    data = read_mapper_output(sys.stdin, separator=separator)
    # groupby groups multiple word-count pairs by word,
    # and creates an iterator that returns consecutive keys and their group:
    #  current_word - string containing a word (the key)
    #  group - iterator yielding all ["<current_word>", "<count>"] items
    for current_word, group in groupby(data, itemgetter(0)):
        try:
            total_count = sum(int(count) for current_word, count in group)
            # output only those words with total_count>5
            if total_count>5:
```

```

        print "%s%s%d" % (current_word, separator, total_count)
    except ValueError:
        # count was not a number, so silently discard this item
        pass
if __name__ == "__main__":
    main()

```

1.b. First 50 words of output file:

```

" And      10
" I        23
" It        7
" My        7
" No,       9
" Oh,      11
" That      6
" The       10
" Tis       12
" Well,     7
" What      6
" Yes,     11
" You       12
" A         46
" Ah        10
" Ah!       8
" Ah,       33
" All       19
" An        7
" And      182
" Are        9
" As        29
" At        17
" Be         9
" Because           12
" But       100
" But,       8
" By        22
" Can        9
" Come      12
" Come,      7
" Defects,"           11
" Did        10
" Do         38
" Dr.        13
" For        27
" Friend     10
" From       10
" Give       6
" Go         13
" God        24
" Good       15
" Great      9
" Ha!        7
" Have       19
" He         82
" Hear,      7
" Here       13
" His        10
" History           7

```

1.c. Last 50 words of output file:

```
"Will 20
"With 20
"Would 8
"Yes 7
"Yes, 127
"Yes," 28
"Yes. 21
"Yes." 27
"Yes; 18
"Yet 7
"You 295
"You're 6
>Your 44
"_I_" 11
"_You_" 12
"a 12
"all 6
"and 109
"as 16
"because 7
"but 67
"by 9
"come 6
"do 10
"for 21
"has 6
"have 6
"he 22
"how 15
"if 30
"in 11
"is 23
"it 30
"it's 7
"let 9
"my 12
"that 100
"the 32
"there 15
"they 7
"this 12
"to 26
"was 6
"we 13
"what 26
"when 9
"who 6
"will 7
"you 49
"your 8
```

## 2.a. Code:

**Mapper Code:** wc\_2\_1\_map.py

```
#!/usr/bin/env python
```

```
import sys
```

```
def read_input(file):
```

```
    for line in file:
```

```
        # split the line into words
```

```
        yield line.strip().split()
```

```
def main(separator='\t'):
```

```
    # input comes from STDIN (standard input)
```

```
    data = read_input(sys.stdin)
```

```
    lastword = ""
```

```
    for words in data:
```

```
        # write each word tab-delimited with the next word and again
```

```
        # tab-delimited with trivial count 1 to STDOUT (standard output)
```

```
        # what we output here will be the input for the
```

```
        # Reduce step, i.e. the input for reducer.py
```

```
        for i in range(len(words)-2):
```

```
            print '%s%s%s%s%d' % (words[i], separator, words[i+1],separator,1)
```

```
        # This is code below deals with producing the bigrams with the word
```

```
        # from current line last word and next line first word
```

```
        if lastword:
```

```
            print '%s%s%s%s%d' % (lastword, separator, words[0],separator,1)
```

```
            lastword=words[-1]
```

```
if __name__ == "__main__":
```

```
    main()
```

**Reducer Code:** wc\_2\_1\_red.py

```
#!/usr/bin/env python
```

```
from itertools import groupby
```

```
from operator import itemgetter
```

```
import sys
```

```
def read_mapper_output(file, separator='\t'):
```

```
    for line in file:
```

```
        yield line.rstrip().split(separator, 2)
```

```
def main(separator='\t'):
```

```
    # input comes from STDIN (standard input)
```

```
    data = read_mapper_output(sys.stdin, separator=separator)
```



```

# groupby groups multiple word-word-count pairs by firstword,
# and creates an iterator that returns consecutive keys and their group:
# current_firstword - string containing a word (the key)
# group - iterator yielding all ["<current_firstword>",
# "<current_secondword>", "<count>"] items
for current_firstword, group_1 in groupby(data, itemgetter(0)):
    try:
        for current_secondword, group_2 in groupby(group_1, itemgetter(1)):
            # groupby groups multiple word-word-count pairs by secondword,
            # and creates an iterator that returns consecutive keys and their group:
            # current_secondword - string containing a word (the key)
            # group - iterator yielding all ["<current_firstword>",
            # "<current_secondword>", "<count>"] items

            try:
                total_count = sum(int(count) for firstword, secondword, count in group_2)
                if total_count > 5:
                    print "%s%s%s%s%d" % (current_firstword, separator,
current_secondword, separator, total_count)
            except ValueError:
                pass
        except ValueError:
            # count was not a number, so silently discard this item
            pass

if __name__ == "__main__":
    main()

```

2.b. First 50 results of output file:

```
"Defects,"      such      11
"Have you      9
"History      of      7
"Information    about    11
"It is      8
"May I      6
"Paradise      Lost,"    8
"Plain Vanilla 22
"Project      Gutenberg"    33
"The Coroner:      6
$5,000) are      20
'AS-IS' WITH      13
($1 to      19
(801) 596-1887,      19
(a) distribution 21
(and you!)      6
(any work      20
(available with 21
(b) alteration, 21
(c) any      21
(does not      21
(if any)      8
(such as      8
(trademark/copyright) agreement. 20
*      *      437
*      *      11
*      *      7
*      *      47
***** This      22
- You      6
- You      10
- You      8
- You      8
--"Paradise      Lost,"    15
.      .      6
.      .      7
.      .      20
1.A. By      20
1.C below.    19
1.C. The      20
1.D. The      20
1.E. Unless    20
1.E.1 through 20
1.E.1. The     20
1.E.2. If      20
1.E.3. If      20
1.E.4. Do      20
1.E.5. Do      20
1.E.6. You     20
1.E.7. Do      20
```

2.c. Last 50 results of output file:

```
weight of 6
whatsoever. You 27
whatsoever. You 7
which was 7
widest variety 10
widest variety 8
willing to 6
willing to 7
willing to 11
willing to 6
window and 6
wish to 12
wish to 6
wish to 6
wish to 10
wishing to 15
wont to 9
worse for 6
worthy of 7
you!) can 21
you," said 6
By and 8
By and 10
Do you 9
It was 6
Pretty soon 12
'Ah!' said 11
'Alas!' said 6
'Behave well, 6
'No,' said 6
'Not wretch 6
'Off with 8
'Open the 6
'Well,' said 9
'it is 7
"Blame it, 8
"Defects," such 10
"Do you 7
"Have you 7
"Information about 9
"It is 7
"It is 6
"MY DEAR 6
"Plain Vanilla 17
"Project Gutenberg" 31
"Queequeg," said 7
"That is 7
"Upon my 7
"Well," said 7
Will you 9
```

### 3.a Code:

The code is same as the one for that is discussed in 3.a. with cleaned data.

### 3.b. Top 100 Results:

```
99056 1 the
57935 1 and
51995 1 of
47126 1 to
32967 1 a
27706 1 in
23114 1 I
21442 1 that
18068 1 his
17738 1 was
16735 1 he
13993 1 with
13747 1 is
13621 1 it
12661 1 as
11918 1 for
10770 1 not
10232 1 be
9954 1 had
9732 1 you
9152 1 by
8815 1 at
8720 1 but
8604 1 on
8233 1 The
8094 1 all
7895 1 my
7798 1 have
7504 1 from
7357 1 or
7132 1 her
7110 1 they
6678 1 which
6467 1 this
6022 1 so
5593 1 she
5573 1 him
5481 1 their
5404 1 are
5175 1 were
4916 1 no
4743 1 And
4733 1 one
4578 1 when
4519 1 we
4264 1 me
4117 1 said
3991 1 an
3961 1 would
3914 1 there
```

4578	1	when	
4519	1	we	
4264	1	me	
4117	1	said	
3991	1	an	
3961	1	would	
3914	1	there	
3896	1	if	
3647	1	been	
3551	1	out	
3525	1	any	
3498	1	what	
3393	1	will	
3291	1	up	
3286	1	into	
3265	1	could	
3255	1	He	
3235	1	some	
3226	1	them	
3105	1	But	
3104	1	your	
3010	1	who	
2943	1	more	
2751	1	very	
2720	1	do	
2646	1	other	
2607	1	upon	
2587	1	then	
2586	1	such	
2531	1	may	
2486	1	than	
2397	1	shall	
2391	1	our	
2379	1	like	
2372	1	can	
2295	1	must	
2286	1	It	
2278	1	man	
2263	1	has	
2203	1	little	
2197	1	see	
2178	1	about	
2122	2	down	Mr.
2099	1	only	
2064	1	should	
2032	1	did	
1997	1	now	
1968	1	before	